

# CSE 5203

## Studio: PPT and practice with binaries

### Overview

Today we will explore a practical example of a security vulnerability that involves the three components of People, Process, and Technology.

Keep detailed notes of your answers below (place your comments and screen shots in between the provided horizontal lines); you may want to refer to these when working on future assignments.

Pay attention to formatting issues when copying lines in this doc to the terminal. You may have to re-type single-quote characters such as ''.

### Part 1: People

For this part, you will be working on your SEED Labs VM, so start that machine now. If you don't have access to a working SEED Labs VM, then you will need to either pause here to get one up and running, or work in a group with someone who has the VM working and use their machine together.

#### 1.1 date

In this exercise, we will be working with the Linux utility program called **date**. Enter the following command lines in a terminal window.

```
date  
date --date="4 hours ago"  
date --date="2 years ago" +%Y%m%d
```

Where is the **date** program in the filesystem? Show how you found it below. Take a screen shot of the terminal window that includes your successful command line(s).

```
[[01/28/26]seed@Kaiyuan:~$ date
Wed Jan 28 16:57:00 UTC 2026
[[01/28/26]seed@Kaiyuan:~$ date --date="4 hours ago"
Wed Jan 28 12:57:09 UTC 2026
[[01/28/26]seed@Kaiyuan:~$ date --date="2 years ago" +%y%m%d
240128
[[01/28/26]seed@Kaiyuan:~$ which date
/usr/bin/date
[[01/28/26]seed@Kaiyuan:~$ ls -l /usr/bin/date
-rwxr-xr-x 1 root root 108920 Sep 5 2019 /usr/bin/date
[01/28/26]seed@Kaiyuan:~$ █
```

Briefly describe what the program does.

The date program displays the current system date and time. It can also format the displayed time based on the options.

## 1.1 my\_date

Download my\_date.zip from Canvas to your VM, extract and run it. Take a few minutes to explore it. (If you have an EIT-provided AWS SEED VM, then you will need to use the scp utility on your terminal command line to copy my\_date.zip from your local laptop to your remote VM.)

Briefly describe what the program does.

Display the current system date and time. It can also format the displayed time based on the options.

Are the two programs similar? Identical? Briefly explain your answer. (For now, don't use diff tools. Simply try to examine the programs as a user.)

---

It's similar. When I use the command "./my\_date", it prints "Wed Jan 28 17:02:21 UTC 2026" which is the same as "date". And it also prints the same result with command "date --date="4 hours ago"" and "date --date="2 years ago" +%Y%m%d"

---

## Part 2: Process

How can you, as a user, protect yourself from downloading malicious files? Any initial thoughts?

---

I only download software from trusted and official sources. I should be cautious of unexpected downloads, email attachments or links from unknown senders. Keeping my operating system up to date can also help reduce the risk of malware.

---

Go over these three articles, published between 11/18-11/20 2019:

- November 18, 2019: <https://github.com/monero-project/monero/issues/6151>
- November 19, 2019: <https://bartblaze.blogspot.com/2019/11/monero-project-compromised.html>
- November 20, 2019: <https://thehackernews.com/2019/11/hacking-monero-cryptocurrency.html>

Would you answer the previous question differently now? Please explain.

---

Yes, it will be different. Now I should verify software integrity by checking cryptographic hashes or digital signatures and prefer building software from source. Articles highlights that protecting against malicious downloads requires both trusting reputable sources and verifying that the downloaded files have not been tampered with.

---

## Part 3: Technology

We can use file signatures and checksums to validate the integrity of the files and programs we download.

A **checksum** is “a small-sized block of data derived from another block of digital data for the purpose of detecting errors that may have been introduced during its transmission or storage.”<sup>1</sup> In other words, checksums help preserve data integrity by detecting whether data has been modified.

A hash function often serves as a checksum for digital data, and “secure” or “cryptographic” hash functions have been designed to be robust to intentional data modification (e.g. by an attacker) as well as benign errors in transmission or storage.

One common secure hash function is the SHA-256 algorithm, which you can read more about [here](#) if you are interested.

Run the following command on both the `date` and `my_date` program files:

```
sha256sum <program_name>
```

Paste your results here (show command and output in a screen shot):

---

---

```
[01/28/26]seed@Kaiyuan:~$ sha256sum /usr/bin/date
a1cb3238a537c39e58c0f25c41e75d22d82b9e92d433185cd00438ba5e8e25a5  /usr/bin/date
[01/28/26]seed@Kaiyuan:~$ ls
Desktop  my_date  my_date.zip
[01/28/26]seed@Kaiyuan:~$ sha256sum my_date
930103a8b531f05143d18efc44c69c8734d1cb1234de1dd9e8752dc746986546  my_date
[01/28/26]seed@Kaiyuan:~$ █
```

---

---

Now, use 'vimdiff' and 'strings' to find the differences between the two files:

```
vimdiff <(strings /path/to/normal/date) <(strings my_date)
```

---

<sup>1</sup> <https://en.wikipedia.org/wiki/Checksum>

Scroll all the way down to near the end of the files, where the text strings are. What are the differences you see in the program text strings? Why would an attacker introduce those differences?

---

The normal date binary contains many detailed error messages and strings related to date parsing, time zones and input validation, such as warnings and parsing errors. In contrast, my\_date is missing most of these descriptive strings.

An attacker may hide malicious functionality in the file to minimize error output, or make the program's bad behavior harder to analyze. Also, without input validation, malicious actions can be more easily injected into the program.

---

Now, let's think like an attacker. Can you modify the `my_date` program to use another domain (e.g. ".com" instead of ".net")? If yes, briefly describe below what you did to modify the program. If you don't know how to do that, move on to the next section, where we'll give you some hints and introduce you to a few helpful tools.

---

I have no idea now.

---

## Part 4: Let's get our hands dirty!

Let's start with `objdump`. Briefly describe what the `objdump` program does.

---

'Objdump' is a tool used to display information about object files and executables.

---

Enter the following command at the terminal, replacing `{loc}` with the location you found above for the built-in `date` program (you probably want

|  |  |  |
|--|--|--|
|  |  |  |
|--|--|--|

to find-replace {loc} for your location so that you can copy and paste directly). This will pipe the output of objdump -d into a program called less. Here are some quick and easy commands to learn to make less easier to navigate. Note that, much like Vi(m) in “command mode,” less interprets your keypresses as commands at the output window.

Vi(m)/less commands: Scroll up (J), scroll down (K), start a forwards search for a string (/), start a backward search for a string (?), quit (q), scroll down one window (SPACE), scroll forwards ½ window (d), and scroll back ½ window (b). Now you know enough to be quick using less!

Explain below what you see when you enter the following command, using only a few sentences.

```
objdump -d {loc}/date | less
```

---

I see the disassembled machine code of the date program. The output shows memory addresses, assembly instructions and function labels that represent the internal structure of the executable. Using ‘less’ allows me to scroll through and search the disassembly more easily.

---

---

Enter the following command at the terminal, contrast this output with the previous output, and explain what the -xtrds switches mean.

```
objdump -xtrds {loc}/date | less
```

---

'objdump -d 'output mainly shows the disassembled machine instructions. And this command displays additional structural information about the binary. The output includes details like file headers, symbol tables, relocation entries and section information.

---

Enter the following command at the terminal, and include the screen shot output below.

---

---

```
objdump -xtrds {loc}/date | grep -A 2 bugs
-----
[01/28/26]seed@Kaiyuan:~$ objdump -xtrds /usr/bin/date | grep -A2 bugs
grep: bugs: No such file or directory
[01/28/26]seed@Kaiyuan:~$
```

Make a copy of `date` in your home directory. Launch either `ghex` or `bless`. Briefly describe the tool you chose and its purpose below. (If you have an EIT-provided AWS SEED VM, skip the remainder of this part and proceed to Part 5.)

Open your local copy of `date` using `bless` or `ghex`, and search for the text 'west'. In the right-hand column, you should see the string "west", and the hex numbers "77 65". Change these numbers to "65 61". Save and quit the file. Go back to the terminal and run (the local version of) "./date --help". Do you see the differences? What is it? What does "65 61" represent?

Now go back to the "attacker" question at the end of part 3, and complete the question if you were not able to before now.

## Part 5: Debug!

In this section, we will practice an important skill: debugging!

## 5.1 Readelf

First, let's look at a useful tool: `readelf`. ELF is the object file format used in Linux and many other systems. Execute the following commands at the terminal, and explain what each one shows you. If you'd like, you can keep a copy of each output in this file for future reference. For this Part, `{loc}` now refers to the location of the `my_date` program on your system.

```
readelf -l {loc}/my_date
readelf -S {loc}/my_date
readelf -W -s {loc}/my_date
readelf -x 16 {loc}/my_date
```

|                       |  |   |
|-----------------------|--|---|
| readelf -l<br>my_date | displays the program headers, showing how the executable is mapped into memory at runtime. | <pre>[01/28/26]seed@Kaiyuan:~\$ readelf -l my_date Elf file type is EXEC (Executable file) Entry point 0x8049b9 There are 9 program headers, starting at offset 52  Program Headers:   Type          Offset   VirtAddr   PhysAddr   FileSiz MemSiz Flg Align   PHDR          0x000034 0x00040034 0x00040034 0x00120 0x00120 R 0x4   INTERP         0x000154 0x00040154 0x00040154 0x00013 0x00013 R 0x1     [Requesting program interpreter: /lib/ld-linux.so.2]   LOAD           0x000000 0x00040000 0x00040000 0x000299 0x000299 RW 0x1000   LOAD           0x000f80 0x00058f80 0x00058f80 0x000299 0x000299 RW 0x1000   DYNAMIC        0x000ff1a 0x00058f1a 0x00058f1a 0x00008 0x00008 RW 0x4   NOTE           0x000168 0x00040168 0x00040168 0x00044 0x00044 RW 0x4   GNU_EH_FRAME   0x00d370 0x00055370 0x00055370 0x0354 0x00354 R 0x4   GNU_STACK      0x000000 0x00000000 0x00000000 0x00000000 0x00000000 RW 0x10   GNU_RELRO     0x000ff08 0x00058f08 0x00058f08 0x0000f8 0x0000f8 R 0x1  Section to Segment mapping: Segment Sections...   00   01 .interp   02 .interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version   _r .rel.dyn .rel.plt .init .plt.got .text .fini .rodata .eh_frame_hdr .eh_frame   .init_array .fini_array .jcr .dynamic .got   04 .dynamic   05 .note.ABI-tag .note.gnu.build-id   06 .eh_frame_hdr   07   08 .init_array .fini_array .jcr .dynamic .got</pre>  |
| readelf -S<br>my_date | lists the section headers of the binary  | <pre>[01/28/26]seed@Kaiyuan:~\$ readelf -S my_date There are 29 section headers, starting at offset 0x102d4:  Section Headers: [Nr] Name             Type            Addr        Off  Size   ES Flg Lk Inf Al [ 0] NULL             PROGBITS        00000000 00000000 00000000 00 0 0 0 0 0 [ 1] .interp          PROGBITS        00000000 00000000 00000000 00 0 0 1 [ 2] .note.ABI-tag    NOTE           00000000 00000000 00000000 00 0 0 4 [ 3] .note.gnu.build-i NOTE           00000000 00000000 00000000 00 0 0 4 [ 4] .gnu.hash         GNU_HASH        00000000 00000000 00000000 04  A 5 0 4 [ 5] .dynsym          DYNSYM         00000000 00000000 00000000 00  A 6 1 4 [ 6] .dynstr          STRTAB         00000000 00000000 00000000 00  A 0 0 1 [ 7] .gnu.version     VERSYM         00000000 00000000 00000000 02  A 5 0 2 [ 8] .gnu.version_r   VERNEED        00000000 00000000 00000000 00  A 6 1 4 [ 9] .rel.dyn          REL            00000000 00000000 00000000 00  A 5 0 4 [10] .rel.plt          REL            00000000 00000000 00000000 08  AI 5 24 4 [11] .init             PROGBITS        00000000 00000000 00000000 00  AX 0 0 4 [12] .plt              PROGBITS        00000000 00000000 00000000 04  AX 0 0 16 [13] .plt.got          PROGBITS        00000000 00000000 00000000 00  AX 0 0 8 [14] .text             PROGBITS        00000000 00000000 00000000 00  AX 0 0 16 [15] .fini             PROGBITS        00000000 00000000 00000000 00  AX 0 0 4 [16] .rodata            PROGBITS        00000000 00000000 00000000 00  AX 0 0 32 [17] .eh_frame_hdr    PROGBITS        00000000 00000000 00000000 00  A 0 0 4 [18] .eh_frame          PROGBITS        00000000 00000000 00000000 00  A 0 0 4 [19] .init_array        INIT_ARRAY     00000000 00000000 00000000 00  WA 0 0 4 [20] .fini_array        FINI_ARRAY     00000000 00000000 00000000 00  WA 0 0 4 [21] .jcr              PROGBITS        00000000 00000000 00000000 00  WA 0 0 4 [22] .dynamic           DYNAMIC         00000000 00000000 00000000 08  WA 6 0 4 [23] .got              PROGBITS        00000000 00000000 00000000 04  WA 0 0 4 [24] .got.plt          PROGBITS        00000000 00000000 00000000 04  WA 0 0 4 [25] .data              PROGBITS        00000000 00000000 00000000 00  WA 0 0 32 [26] .bss               NOBITS         00000000 00000000 00000000 00  WA 0 0 64 [27] .gnu_debuglink    PROGBITS        00000000 00000000 00000000 00  WA 0 0 1 [28] .shstrtab          STRTAB         00000000 00000000 00000000 00  0 0 1  Key to Flags:  W (write), A (alloc), X (execute), M (merge), S (strings), I (info), </pre> |

|                       |  |  |
|-----------------------|--|--|
| readelf -W -s my_date | displays the symbol table, including function and variable symbols used by the program | <pre>[01/28/26]seed@Kaiyuan:~\$ readelf -W -s my_date Symbol table '.dynsym' contains 84 entries: Num: Name                 Value      Size Type    Bind   Vis      Ndx  0: .ctors               0           0 FUNC    LOCAL  DEFAULT     UND  1: _GLOBAL_OFFSET_TABLE_ 0           0 FUNC    GLOBAL DEFAULT     UND  2: __ctype_toupper_loc@GLIBC_2.3 (2)  3: __unisetenv@GLIBC_2.0 (3)  4: __strcmpl@GLIBC_2.0 (3)  5: __open64@GLIBC_2.1 (4)  6: __fflush64@GLIBC_2.0 (3)  7: __exit@GLIBC_2.0 (3)  8: __free@GLIBC_2.0 (3)  9: __memcpy@GLIBC_2.0 (3) 10: __mbssinit@GLIBC_2.0 (3) 11: __lock_getname@GLIBC_2.17 (5) 12: __flock_getname@GLIBC_2.3 (4) 13: __time@GLIBC_2.0 (3) 14: __fseek64@4@GLIBC_2.1 (4) 15: __memcmp@GLIBC_2.0 (3) 16: __gettimeofday@GLIBC_2.0 (3) 17: __dgettext@GLIBC_2.0 (3) 18: __stack_chk_fail@GLIBC_2.4 (6) 19: __textdomain@GLIBC_2.0 (3) 20: __getline@GLIBC_2.0 (3) 21: __fopen64@4@GLIBC_2.1 (4) 22: __tzset@GLIBC_2.0 (3) 23: __brtowc@GLIBC_2.0 (3) 24: __ctype_get_mb_cur_max@GLIBC_2.0 (3) 25: __fpending@GLIBC_2.3 (7) 26: __mbrtowc@GLIBC_2.0 (3) 27: __cxa_atexit@GLIBC_2.1.3 (8) 28: __error@GLIBC_2.0 (3) 29: __getenv@GLIBC_2.0 (3)</pre>  |
| readelf -x 16 my_date | dumps the raw hexadecimal contents of section number 16                                | <pre>[01/28/26]seed@Kaiyuan:~\$ readelf -x 16 my_date Hex dump of section '.rodata': 0x000521a0 01000200 74696d65 20257320 .....time % 0x000521b0 6973206f 7574206f 66207261 66676500 is out of range. 0x000521c0 5b007465 73742069 66766f63 6174696f [.test invocation 0x000521d0 6e004d75 6c74692d 63616cc6 20696e76 n.Multi-call inv 0x000521e0 6f636174 696f6e00 73686132 32347375 ocation.sha24su 0x000521f0 6d007368 61322875 74696c69 74696573 m.sha2 utilities 0x00052200 00736861 32353673 7560073 68613338 .sha256sum.sha38 0x00052210 3473756d 00736861 35313273 756d000a 4sum.sha512sum.. 0x00052220 2573206f 6e6c696e 65206865 6c703a20 %s online help; 0x00052230 3c25733e 0a00474e 5520636f 72657574 &lt;%s&gt;..GNU coreut 0x00052240 696c7300 656e5f00 25612025 62202565 ills.en.%a %b %e 0x00052250 2025483a 254d3a25 53202558 20255900 %H:%M:%S %Z %Y. 0x00052260 2f757372 2f736861 72652f6 6f63616c /usr/share/local 0x00052270 65002d2d 7266632d 33333339 002d2d69 e.-rfc-3339.--i 0x00052280 736f2d38 36303100 545a3d55 54433000 so-8601.TZ=UTC0. 0x00052290 44617669 64204d61 634b656e 7a696500 David Mackenzie. 0x000522a0 65787472 61206f70 6572616a 64202573 extra operand %s 0x000522b0 00545a00 7374616e 64617264 20696e70 .TZ.standard inp 0x000522c0 75740069 0e76616c 69642064 61746520 ut.invalid date 0x000522d0 25730063 616e66ef 74207365 74206461 %s.cannot set da 0x000522e0 74650066 696c6500 72656665 72656e63 te.file.referenc 0x000522f0 65007266 632d3832 32007266 632d3238 e.rfc-822.rfc-28 0x00052300 32320073 65740075 63740075 74630075 22.set.utc.u 0x00052310 6e697665 7273616c 00686564 70007665 niversal.help.ve 0x00052320 7273696f 6e006861 75727300 6d696e75 rision.hours.minu 0x00052330 74657300 7365636f 6e647300 6e730000 tes.seconds.ns.. 0x00052340 54727920 27257320 2d2d6865 6c702720 Try '%s --help' 0x00052350 666f7220 6d6f7265 20696e66 6f726d61 for more informa 0x00052360 74696f6e 2e0a0000 55736167 653a2025 tion....Usage: %</pre> |

## 5.2 Using gdb

At the terminal, enter the following  
gdb  
(gdb) help

Use the interactive help feature to explore gdb's options.

## 5.3 Examining processes

At the terminal, enter the following.  
gdb my\_date  
(gdb) run

```
(gdb) quit
```

Copy your output below. Copy-paste the text rather than a screen shot.

---

```
[01/28/26]seed@Kaiyuan:~$ gdb my_date
GNU gdb (Ubuntu 9.2-0ubuntu1~20.04.2) 9.2
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later
<http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
```

For help, type "help".

```
Type "apropos word" to search for commands related to "word"...
/opt/gdbpeda/lib/shellcode.py:24: SyntaxWarning: "is" with a literal.
Did you mean "=="?
    if sys.version_info.major is 3:
/opt/gdbpeda/lib/shellcode.py:379: SyntaxWarning: "is" with a literal.
Did you mean "=="?
    if pyversion is 3:
Reading symbols from my_date...
(No debugging symbols found in my_date)
gdb-peda$ run
Starting program: /home/seed/my_date
Wed Jan 28 18:19:44 UTC 2026
[Inferior 1 (process 41066) exited normally]
Warning: not running
gdb-peda$ quit
[01/28/26]seed@Kaiyuan:~$
```

---

Try the following commands to get a feel for the tool.

```
gdb my_date
(gdb) set args "--help"
(gdb) break __libc_start_main
(gdb) run
(gdb) frame
(gdb) bt
(gdb) info frame
(gdb) info registers
```

```
(gdb) x /16xw $esp
```

Explore other registers and memory locations. What can you say about where the code and stack are located?

---

In GDB, the instruction pointer points into the executable code region, showing where the current code is executing. The stack pointer points into the stack region, which stores return addresses and saved data for function calls. The stack and code addresses are in a different region. And the stack addresses typically grow toward lower addresses as functions are called. This separation reflects that code and stack occupy different mapped regions with different permissions.

---

Now try the following.

```
gdb my_date
(gdb) set args "--version"
(gdb) break __libc_start_main
(gdb) run
(gdb) maintenance info sections
0x80521a0->0x805536f at 0x0000a1a0: .rodata ALLOC LOAD READONLY DATA
HAS_CONTENTS
note the start address of the .rodata section
(gdb) x/20s {.rodata address}
note the start address of the string "GNU coreutils" as {start}
(gdb) set *{start} = 0x20554c42
(gdb) c
```

What does the code above do? Can you modify it to do something else that demonstrates that you understand how it works? Did you run into a snag? Provide your explanation and your modification below. A reference like this might be helpful: <http://www.asciitable.com>

---

It set a breakpoint at the beginning of the function to pause the program before main() run. The command "maintenance info sections" locates the .rodata section in memory. I used an extra command "searchmem "GNU" {start} {end}" to find the start address of constant string "GNU coreutils". Next, I used the command "x/20s" to confirm I found the correct start address. Then overwrite the first four bytes of the string by using "set". Setting \*{start} = 0x20554c42 changes the bytes to 42

4c 55 20, which is "BLU ". The final c continues the program and the output is the modified string "BLU coreutils" .

I also modified the name "David MacKenzie" to "Kaiyuan Xu" through the same method.

```
[gdb-peda]$ searchmem "David" 0x80521a0 0x805536f
Searching for 'David' in range: 0x80521a0 - 0x805536f
Found 1 results, display max 1 items:
my_date : 0x8052290 ("David MacKenzie")
[gdb-peda]$ set {char[15]} 0x8052290 = "Kaiyuan Xu"
[gdb-peda]$ c
Continuing.
date (GNU coreutils) 8.25
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.net/licenses/gpl.html>.
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.

Written by Kaiyuan Xu.
[Inferior 1 (process 42408) exited normally]
Warning: not running
gdb-peda$
```

Snag: Before learning the command 'searchmem', I'm confused how to find the start address of the specific words. I wondered if I should count the bites and modify the content after "GNU coreutils". By searching among the websites, I finally found this command which helped me much easier to locate the first address of the word.

---

## COMPLETE

That's all! Please submit a .pdf version of this file to Gradescope.